The Ω-Gradient: Evidence-Weighted Multi-Agent Policy Optimization with Opponent Shaping, Cooperative Communication, and the Blessing of Dimensionality

Technical Report

Eugene Shcherbinin
London School of Economics
y.shcherbinin@lse.ac.uk

March 2026

Abstract

We present a unified framework for multi-agent policy gradient methods in general stochastic games, building on the convergence theory of Giannou et al. (2022). Our framework introduces four orthogonal improvements to the standard policy gradient, each addressing a distinct source of suboptimality: **(1)** Evidence-weighted updates (EW-PG) reduce gradient variance by a factor of $\mathrm{HM}(V)/\,\mathrm{AM}(V) \leq 1$ using Keynesian evidence weights; **(2)** Opponent-shaping updates (LOLA-PG) enlarge the basin of attraction via spectral reinforcement of the gradient field's Jacobian; **(3)** Cooperative communication (Coop-PG) enables coalition formation through lossy policy transmission, with quality bounded by agents' self-knowledge (the $\mathcal{O} \to \Pi$ inference gap); **(4)** Sparse policy regularization exploits the blessing of dimensionality to achieve $O(k \log d)$ sample complexity in d-dimensional action spaces with k-sparse optima. All four mechanisms compose into a single update rule — the Ω-gradient — with provable convergence to stable Nash equilibria. We provide Lean 4 formalizations of the core theorems and validate all predictions through numerical experiments on matrix games.

Contents

1 Introduction **4**
1.1 Contributions 4
1.2 Relation to Prior Work 5

2 Preliminaries **5**
2.1 Stochastic Games 5
2.2 Policy Gradient and REINFORCE 5
2.3 The SOS Condition (Giannou et al.) 5

3 Evidence-Weighted Policy Gradient **6**
3.1 Motivation: Heterogeneous Information 6
3.2 The Keynesian Evidence Weight 6
3.3 Variance Improvement 6
3.4 Algorithm 7

4 Opponent-Shaping Policy Gradient **7**
4.1 The LOLA Correction . . . 7
4.2 Convergence Results . . . 7

5 Cooperative PG with Lossy Communication **7**
5.1 The Self-Knowledge Problem . . . 7
5.2 Communication Loss Decomposition . . . 8
5.3 Triple-Duty Evidence Weight . . . 8
5.4 Coalition Formation . . . 8

6 The Blessing of Dimensionality **8**
6.1 The Apparent Curse . . . 8
6.2 The Blessing: Sparse Recovery . . . 9
6.3 The Natural Curriculum . . . 9

7 Fixed-Point Nash Equilibrium Search **9**
7.1 Nash Equilibria as Fixed Points . . . 9
7.2 Contraction and the Banach Fixed-Point Theorem . . . 10
7.3 Bayesian Exploration of the Fixed-Point Set . . . 10
7.4 Cooperation from Selfish Search . . . 10
7.5 The Topological Cooperation Theorem . . . 11
7.6 Integration: FP-NE as Phase 0 of Ω-PG . . . 11
7.7 Relation to Meta-Value Learning and Meta-MAPG . . . 12

8 The Complete Ω-Gradient **12**
8.1 Unified Update Rule . . . 12
8.2 Mapping to the Ω-Framework . . . 13
8.3 Six-Way Risk Decomposition . . . 13

9 Lean 4 Formalization **13**

10 Experimental Validation **14**
10.1 Experiment 1: EW-PG Variance Reduction . . . 14
10.2 Experiment 2: Blessing of Dimensionality . . . 15
10.3 Experiment 3: Coalition Formation . . . 15
10.4 Experiment 4: Full Ω-PG . . . 15
10.5 Experiment 5: Sparsity Curriculum . . . 16
10.6 Experiment 6: FP-NE Discovery and Contraction Rates . . . 16
10.7 Experiment 7: Equilibrium Selection Quality . . . 17
10.8 Experiment 8: Bayesian Stopping and Cooperation from Selfishness . . . 18
10.9 Experiment 9: Iterated Prisoner's Dilemma — Adaptation (Kim et al.) . . . 18
10.10 Experiment 10: Iterated RPS — Competitive Exploitation . . . 19
10.11 Experiment 11: Strategy Identification . . . 19
10.12 Experiment 12: Scaling to N Players . . . 19
10.13 Experiment 13: Topological Cooperation Theorem . . . 20
10.14 Experiment 14: AI Safety — Corrigibility . . . 20
10.15 Experiment 15: AI Safety — Deceptive Alignment . . . 20
10.16 Experiment 16: AI Safety — Debate . . . 20
10.17 Experiment 17: Mesa-Optimisation Detection . . . 21
10.18 Experiment 18: Power-Seeking Dynamics . . . 22
10.19 Experiment 19: Multi-Principal AI Governance . . . 22
10.20 Experiment 20: The Alignment Tax . . . 23

10.21 Summary of Experimental Findings 25

11 Discussion **26**
11.1 The Evidence Weight as Unifying Concept 26
11.2 LOLA and Cooperation as Dual Gödelian Mechanisms 26
11.3 The Vibing Scale: A Horseshoe, Not a Line 26
11.4 Computation as Cooperation Mechanism: Implications for AI Safety 28
11.5 Limitations and Future Work 28

12 Conclusion **29**

1 Introduction

Multi-agent reinforcement learning (MARL) in general-sum stochastic games faces three fundamental challenges beyond single-agent RL:

1. **Heterogeneous information**: Agents have different observation histories, leading to unequal gradient quality. Standard policy gradient (PG) treats all agents identically, wasting the structure in this heterogeneity.

2. **Strategic interaction**: Agents' rewards depend on others' policies. Ignoring this coupling (independent learning) can lead to cycling, divergence, or convergence to suboptimal equilibria with small basins of attraction.

3. **Scalability**: As action spaces grow, REINFORCE variance scales linearly in the number of actions d, rendering naive PG impractical for combinatorial or continuous action spaces.

We address all three challenges through a unified framework rooted in the Ω-framework's five-component gradient decomposition:

$$\boxed{\pi_{i,n+1} = \mathrm{proj}_{\Pi}\left(\pi_{i,n} + \gamma_n \cdot \underbrace{w_{i,n}}_{\text{evidence}} \cdot \left(\underbrace{\hat{v}_{i,n}}_{\text{PG}} + \underbrace{\lambda_n \cdot \mathrm{OS}_{i,n}}_{\text{opponent shaping}} + \underbrace{\beta_n \cdot \mathrm{coop}_{i,n}}_{\text{cooperation}}\right)\right)} \tag{1}$$

with L1 proximal regularization for sparsity when d is large.

1.1 Contributions

1. **Evidence-Weighted PG** (§3): Agents weight updates by $w_i = V_{\min}/V_i$, where V_i is the Keynesian evidence variance. Variance improves by $\mathrm{HM}(V)/\,\mathrm{AM}(V)$. Convergence rate preserved at $O(1/n^q)$.

2. **Opponent-Shaping PG** (§4): LOLA-type corrections with the first convergence guarantee in general stochastic games. Annealed shaping preserves rate; constant shaping enlarges the basin of attraction under a spectral reinforcement condition.

3. **Cooperative PG with Lossy Communication** (§5): Coalition formation through policy communication. Self-knowledge bottleneck: communication quality $\leq$ evidence quality. The evidence weight does triple duty (gradient, self-knowledge, communication).

4. **Blessing of Dimensionality** (§6): For k-sparse optimal policies in d-dimensional action spaces, L1-regularized EW-PG achieves $O(k \log(d/k))$ sample complexity via concentration of measure and sparse recovery.

5. **Fixed-Point NE Search** (§7): A meta-algorithm that explicitly discovers Nash equilibria as fixed points of the smoothed best-response map BR_τ (Banach contraction), with Bayesian stopping via Good-Turing/Chao1 estimators. All selfish agents benefit from thorough search—cooperation emerges from the topology of the fixed-point set.

6. **Lean 4 Formalization** (§9): Core theorems mechanically verified, including the energy inequality, AM-HM bound, basin enlargement, and communication loss decomposition.

7. **Experimental Validation** (§10): Eight experiments on matrix games and PettingZoo environments confirming all theoretical predictions, including FP-NE discovery, Bayesian stopping, and equilibrium selection quality.

1.2 Relation to Prior Work

Convergence theory. Giannou et al. (2022) proved the first convergence results for PG in general stochastic games under a Stable Opponent-Sensitive (SOS) condition. Our work extends their framework in four orthogonal directions while preserving their convergence guarantees.

Opponent shaping. Foerster et al. (2018) introduced LOLA but provided no convergence guarantees beyond bimatrix games. Letcher et al. (2019) analyzed stability but not convergence rates. We provide the first rate-preserving convergence theorem for LOLA-type methods in general stochastic games.

Communication in MARL. CommNet (Sukhbaatar et al., 2016), TarMAC (Das et al., 2019), and DIAL (Foerster et al., 2016) learn communication protocols but lack convergence theory. Our framework provides information-theoretic bounds on communication quality tied to evidence.

High-dimensional RL. The curse of dimensionality in RL is well-studied. Our contribution connects compressed sensing and concentration of measure to policy gradient variance, showing that sparsity transforms the curse into a blessing.

Nash equilibrium computation. Nash's existence theorem (1950) via Brouwer/Kakutani guarantees at least one mixed NE in finite games. The Lemke-Howson algorithm (1964) finds one NE in bimatrix games; support enumeration finds all. McKelvey and Palfrey (1995) introduced quantal response equilibria (QRE) as smoothed fixed points of the logit best-response map. Our FP-NE search combines QRE computation with Bayesian species-discovery (Good-Turing/Chao1 estimators) for the first explore-exploit framework over the NE set.

Meta-learning in MARL. MeVa (Cooijmans et al., 2023) learns meta-value functions over policy space via implicit Bellman equations. Meta-MAPG (Kim et al., 2021) derives a three-term meta-gradient decomposition accounting for own and peer learning dynamics. Both address non-stationarity but neither provides Bayesian guarantees on equilibrium discovery.

2 Preliminaries

2.1 Stochastic Games

An N-player stochastic game $\mathcal{G} = (N, \mathcal{S}, \{\mathcal{A}_i\}, \{R_i\}, P, \gamma)$ consists of:

- N agents with action spaces $\mathcal{A}_i$, $|\mathcal{A}_i| = d_i$
- State space $\mathcal{S}$, transition P, discount γ
- Reward functions $R_i : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$

Each agent i has a policy $\pi_i \in \Pi_i = \Delta(\mathcal{A}_i)^{|\mathcal{S}|}$, and the joint policy space is $\Pi = \prod_i \Pi_i \subset \mathbb{R}^d$ with $d = \sum_i d_i \cdot |\mathcal{S}|$.

2.2 Policy Gradient and REINFORCE

The policy gradient for agent i is $v_i(\pi) = \nabla_{\pi_i} J_i(\pi)$. The REINFORCE estimator:

$$\hat{v}_i(\pi) = R(\tau) \cdot \nabla_{\pi_i} \log \pi_i(\tau) \tag{2}$$

is unbiased ($\mathbb{E}[\hat{v}_i] = v_i$) with variance σ_i^2.

2.3 The SOS Condition (Giannou et al.)

Definition 2.1 (Stable Opponent-Sensitive Nash)**.** *π^* is an SOS Nash with parameter $\mu > 0$ if there exists $\rho > 0$ such that for all π with $\|\pi - \pi^*\| < \rho$:*

$$\langle v(\pi), \pi - \pi^* \rangle \leq -\mu \|\pi - \pi^*\|^2. \tag{3}$$

This requires the gradient field to point "inward" near π^* — a contraction condition stronger than Nash but weaker than strict Nash. Giannou et al. show that under SOS, PG converges almost surely with rate $\mathbb{E}[\|\pi_n - \pi^*\|^2] = O(C/n^q)$.

3 Evidence-Weighted Policy Gradient

3.1 Motivation: Heterogeneous Information

In multi-agent settings, agents have unequal information quality. Agent i's gradient variance $V_i = \mathrm{Var}(\hat{v}_i)$ reflects its *evidence* — the informativeness of its observation history. High-variance agents (low evidence) inject noise into the joint dynamics, slowing convergence for everyone.

3.2 The Keynesian Evidence Weight

Inspired by Keynes's logical probability and the Ω-framework's truth function, we define:

$$w_{i,n} = \frac{V_{\min,n}}{V_{i,n}}, \qquad V_{\min,n} = \min_j V_{j,n}. \tag{4}$$

The EW-PG update replaces γ_n with $\gamma_n \cdot w_{i,n}$:

$$\pi_{i,n+1} = \mathrm{proj}_\Pi\big(\pi_{i,n} + \gamma_n \cdot w_{i,n} \cdot \hat{v}_{i,n}\big). \tag{5}$$

3.3 Variance Improvement

Theorem 3.1 (AM-HM variance improvement)**.** *Under EW-PG with weights* (4)*, the effective variance of the joint gradient satisfies:*

$$\sigma_w^2 = \frac{\mathrm{HM}(V)}{\mathrm{AM}(V)} \cdot \sigma_{std}^2 \le \sigma_{std}^2 \tag{6}$$

where $\mathrm{AM}(V) = \frac{1}{N}\sum_i V_i$ *and* $\mathrm{HM}(V) = N/\sum_i V_i^{-1}$*. Equality holds iff all* V_i *are equal.*

Proof. The weighted variance sum is $\sum_i \gamma_n^2 w_i^2 \sigma_i^2$. With $w_i = V_{\min}/V_i$ and $\sigma_i^2 \propto V_i$: $\sum_i w_i^2 V_i = V_{\min}^2 \sum_i V_i^{-1}$. The standard sum is $\sum_i V_i$. Their ratio: $\frac{V_{\min}^2 \sum_i V_i^{-1}}{\sum_i V_i} = \frac{N \cdot V_{\min}^2 \cdot (N/\mathrm{HM})}{N \cdot \mathrm{AM}} = \frac{\mathrm{HM}}{\mathrm{AM}} \cdot \frac{V_{\min}^2}{\mathrm{HM}^2/N}$. After normalization, the improvement factor is HM / AM ≤ 1 by the AM-HM inequality (a consequence of Cauchy-Schwarz). □

Theorem 3.2 (Convergence of EW-PG)**.** *Under the conditions of Giannou et al. Theorem 1, EW-PG converges to stable Nash with rate:*

$$\mathbb{E}[\|\pi_n - \pi^*\|^2 \mid \mathcal{E}] = O\left(\frac{\mathrm{HM}(V)}{\mathrm{AM}(V)} \cdot \frac{C}{n^q}\right) \tag{7}$$

where $q = \min(\ell_b, p - 2\ell_\sigma)$ *and* C *is Giannou's constant.*

3.4 Algorithm

Algorithm 1: Evidence-Weighted Policy Gradient (EW-PG)

Input: Game $\mathcal{G}$, learning rates $\gamma_n = \gamma/(n+m)^p$, decay α

for $n = 1, 2, \ldots$ **do**
  for $i = 1, \ldots, N$ **do**
    $\hat{v}_{i,n} \leftarrow$ REINFORCE estimate;
    $\hat{V}_{i,n} \leftarrow \alpha \hat{V}_{i,n-1} + (1-\alpha)\|\hat{v}_{i,n} - \bar{v}_i\|^2$;
  $V_{\min,n} \leftarrow \min_i \hat{V}_{i,n}$;
  for $i = 1, \ldots, N$ **do**
    $w_{i,n} \leftarrow V_{\min,n}/\hat{V}_{i,n}$;
    $\pi_{i,n+1} \leftarrow \mathrm{proj}_\Pi(\pi_{i,n} + \gamma_n w_{i,n} \hat{v}_{i,n})$;

4 Opponent-Shaping Policy Gradient

4.1 The LOLA Correction

Foerster et al. (2018) proposed augmenting agent i's gradient with an opponent-shaping (OS) term that anticipates how opponents will respond:

$$\mathrm{OS}_i(\pi) = \sum_{j \neq i} \frac{\partial R_i}{\partial \theta_j} \cdot \frac{\partial \theta'_j}{\partial \theta_i} \tag{8}$$

where $\theta'_j = \theta_j - \eta_j \nabla_{\theta_j} R_j$ is j's anticipated one-step update.

4.2 Convergence Results

Theorem 4.1 (Annealed LOLA preserves convergence)**.** *With shaping schedule $\lambda_n = \lambda/(n+m)^r$ and $r > 1 - p$, the LOLA-PG update $\pi_{i,n+1} = \mathrm{proj}_\Pi(\pi_{i,n} + \gamma_n(\hat{v}_{i,n} + \lambda_n \cdot OS_{i,n}))$ converges to stable Nash at the same rate as standard PG.*

Proof sketch. The OS term adds $\lambda_n G$ to the bias bound. The augmented bias $B_n^{\mathrm{LOLA}} = O(n^{-\min(\ell_b, r)})$ satisfies Giannou's condition $p + \min(\ell_b, r) > 1$ when $r > 1 - p$. □

Theorem 4.2 (Basin enlargement under spectral reinforcement)**.** *Let $H = \nabla_\pi OS(\pi^*)$ be the OS Hessian at Nash, with symmetric part $S_H = (H + H^\top)/2$. If S_H is negative semi-definite with largest eigenvalue $-\mu_H$ ($\mu_H \geq 0$), then LOLA-PG has SOS parameter:*

$$\mu_{LOLA} = \mu + \lambda \mu_H > \mu. \tag{9}$$

The spectral reinforcement condition holds in zero-sum games (where the antisymmetric part of Jac_v dominates) and games with negative cross-derivatives.

5 Cooperative PG with Lossy Communication

5.1 The Self-Knowledge Problem

Agent i's policy π_i is an implicit function of its parameters. The agent does not have direct introspective access; it can only *infer* a self-model $\hat{\pi}_i$ from observations:

Definition 5.1 (Self-knowledge loss)**.** $L_{self}^{(i)} = \mathbb{E}[\|\pi_i - \hat{\pi}_i\|^2]$.

Axiom 5.1 (Self-knowledge bounded by evidence)**.** $L_{self}^{(i)} \leq V_i$.

Agents with high variance ("vibing") don't know their own policy well.

5.2 Communication Loss Decomposition

When agent i communicates to teammate j through a lossy channel:

$$\pi_i \xrightarrow{\text{self-model}} \hat{\pi}_i \xrightarrow{\text{channel}} \tilde{\pi}_i^{(j)} \tag{10}$$

Theorem 5.2 (Loss decomposition)**.**

$$\underbrace{L_{comm}^{(i\to j)}}_{total} \leq 2 \underbrace{L_{self}^{(i)}}_{self\text{-}knowledge} + 2\underbrace{L_{channel}^{(i\to j)}}_{channel} \leq 2V_i + 2L_{channel}. \tag{11}$$

Corollary 5.3 (Self-knowledge bottleneck)**.** *Even with a perfect channel, $L_{comm} \leq 2V_i$. An agent that doesn't know its own policy cannot communicate it.*

5.3 Triple-Duty Evidence Weight

The evidence weight w_i simultaneously measures:

1. **Gradient quality**: $\sigma_w^2 = (w_i)^2\sigma_i^2 \leq \sigma_i^2$
2. **Self-knowledge**: $L_{\text{self}} \leq V_i = V_{\min}/w_i$
3. **Communication**: $L_{\text{comm}} \leq 2V_{\min}/w_i + 2L_{\text{ch}}$

All three are manifestations of the same information content.

5.4 Coalition Formation

Definition 5.4 (Rational coalition)**.** *$S \subseteq \{1, \ldots, N\}$ is rational if $V(S) > \sum_{i\in S} V(\{i\})$.*

This applies in *competitive* games too: temporary alliances form when the coordination premium exceeds individual payoffs.

Theorem 5.5 (Cooperative basin enlargement)**.** *Under cooperative reinforcement (symmetric part of the cooperative Hessian is negative semi-definite), the SOS parameter improves to $\mu_{coop} = \mu + \beta\mu_C$.*

6 The Blessing of Dimensionality

6.1 The Apparent Curse

REINFORCE variance decomposes by action:

$$\text{Var}(\hat{g}) = \sum_{a=1}^{d} \pi(a) \cdot \left(\frac{\nabla\pi(a)}{\pi(a)}\right)^2 \cdot R(a)^2 - \|\nabla J\|^2. \tag{12}$$

For low-probability action a with $\pi(a) = \epsilon$, the contribution scales as $R(a)^2/\epsilon \to \infty$. The apparent conclusion: variance is $\Theta(d)$, so REINFORCE fails for large d.

6.2 The Blessing: Sparse Recovery

When the optimal policy is k-sparse ($k \ll d$ actions with positive probability), high-dimensional phenomena reverse the curse.

Definition 6.1 (Effective dimension)**.** $d_{\mathit{eff}} = \mathrm{tr}(\Sigma)/\lambda_{\max}(\Sigma)$ *where* $\Sigma = \mathit{Cov}(\hat{g})$.

Theorem 6.2 (Sub-linear sample complexity)**.** *For a k-sparse optimal policy in a d-action game, the gradient can be estimated to accuracy ϵ with:*

$$M = O\left(\frac{k \log(d/k)}{\epsilon^2}\right) \tag{13}$$

samples, sub-linear in d.

This follows from the restricted isometry property (Candès–Tao, 2005) applied to the REINFORCE gradient measurement ensemble.

6.3 The Natural Curriculum

L1 regularization + evidence weighting creates self-organizing phases:

Phase	μ_n (L1)	d_{eff}	**Behavior**
Early	Large	$\approx k$	Sparse, fast convergence
Middle	Decreasing	k–$2k$	Expand support, calibrated
Late	≈ 0	$\leq d$	Full flexibility, fine-tune

Theorem 6.3 (Convergence of Sparse-EW-PG)**.** *With k-sparse Nash, L1 schedule $\mu_n = \mu_0/(1+n/n_0)$, and effective evidence weights, the convergence constant is:*

$$C_k = O(k \log(d/k)) \quad \textit{(sub-linear in } d\textit{)}. \tag{14}$$

7 Fixed-Point Nash Equilibrium Search

The four mechanisms above (EW, LOLA, Coop, Sparsity) improve the Ω-gradient's *convergence from a given initialization.* But a fundamental problem remains: in games with multiple Nash equilibria, the choice of initialization determines *which* equilibrium the gradient converges to. This section introduces a meta-algorithm that explicitly searches the fixed-point landscape before running Ω-PG, ensuring convergence to the best discoverable equilibrium.

7.1 Nash Equilibria as Fixed Points

A Nash equilibrium is a fixed point of the joint best-response correspondence $\mathrm{BR} : \Delta \rightrightarrows \Delta$, where $\Delta = \prod_i \Delta(\mathcal{A}_i)$ is the product of strategy simplices. The existence of at least one such fixed point is guaranteed by Kakutani's theorem [14] (for the set-valued BR) or equivalently by Nash's original argument via Brouwer's fixed-point theorem [13].

Theorem 7.1 (Nash 1950, via Brouwer/Kakutani)**.** *Every finite N-player game has at least one mixed-strategy Nash equilibrium.*

The proof uses two key structural facts:

1. Δ is compact and convex (finite-dimensional product of simplices).
2. BR is an upper hemicontinuous correspondence with convex, nonempty values.

Kakutani's theorem then guarantees a fixed point $\pi^* \in \mathrm{BR}(\pi^*)$.

For computation, we smooth the (set-valued) BR into a continuous *softmax best response*:

$$\mathrm{BR}_{\tau,i}(\pi_{-i})(a) = \frac{\exp\big(\tau^{-1} \cdot V_i(a, \pi_{-i})\big)}{\sum_{a'} \exp\big(\tau^{-1} \cdot V_i(a', \pi_{-i})\big)} \tag{15}$$

where $V_i(a, \pi_{-i}) = \sum_{a_{-i}} \pi_{-i}(a_{-i}) R_i(a, a_{-i})$. Fixed points of BR_τ are *quantal response equilibria* (QRE; McKelvey & Palfrey, 1995), which converge to NE as $\tau \to 0$.

7.2 Contraction and the Banach Fixed-Point Theorem

Theorem 7.2 (Contraction of BR_τ)**.** *For any finite game with $\|R_i\|_\infty \le M$, the smoothed best-response map $\mathrm{BR}_\tau : \Delta \to \Delta$ satisfies*

$$\|\,\mathrm{BR}_\tau(\pi) - \mathrm{BR}_\tau(\pi')\|_\infty \le \frac{M}{\tau} \cdot \|\pi - \pi'\|_\infty. \tag{16}$$

When $\tau > M$, BR_τ is a contraction mapping with Lipschitz constant $c = M/\tau < 1$, and by Banach's fixed-point theorem it has a unique fixed point reached by iteration at geometric rate c^t.

Proof. The softmax $\sigma_\tau(x)$ has Lipschitz constant $1/(2\tau)$ in $\|\cdot\|_\infty$. The linear payoff map $\pi_{-i} \mapsto V_i(\cdot, \pi_{-i})$ has operator norm $\le M$. By composition, $\|\,\mathrm{BR}_{\tau,i}(\pi) - \mathrm{BR}_{\tau,i}(\pi')\| \le (M/\tau)\|\pi - \pi'\|$. □

Remark 7.3. *When $\tau \le M$, multiple fixed points may exist—this is exactly the regime where equilibrium selection matters and where our Bayesian search adds value. Decreasing τ from $> M$ toward 0 traces the* QRE correspondence*, and bifurcation points correspond to equilibrium selection events.*

7.3 Bayesian Exploration of the Fixed-Point Set

We frame the problem of finding the *best* NE as a species-discovery problem. Each "search" samples a random initialization, runs damped BR_τ iteration to convergence, and records the resulting fixed point. After n searches discovering d distinct equilibria, we estimate:

Total NE count via the Chao1 estimator [16]:

$$\hat{K} = d + \frac{f_1^2}{2 f_2} \tag{17}$$

where f_k = number of NE discovered exactly k times.

Stopping criterion: For each agent i, the *exploration value*

$$\mathrm{EV}_i = \mathbb{P}(\exists \text{ undiscovered } \pi^* : V_i(\pi^*) > V_i^{\text{best}}) \le \underbrace{\left(1 - \frac{d}{\hat{K}}\right)}_{\mathbb{P}(\text{more NE})} \times \underbrace{p_{\text{better}}}_{\mathbb{P}(\text{it's better})} \tag{18}$$

Stop when $\mathrm{EV}_i < \delta$ for all agents i.

7.4 Cooperation from Selfish Search

The central insight: *in games with multiple NE, thorough search for fixed points is a public good that selfish agents voluntarily provide.*

Proposition 7.4 (Search as Public Good)**.** *In a game with $K \geq 2$ NE where payoffs across equilibria are positively correlated ($Corr(V_i(\pi^{*,k}), V_j(\pi^{*,k})) > 0$ over k), increasing the search budget benefits all agents:*

$$\frac{\partial}{\partial n}\,\mathbb{E}\left[\max_{k \leq d(n)} V_i(\pi^{*,k})\right] > 0 \quad \forall i.$$

This is a functional-analytic grounding of cooperation: it arises not from altruism or mechanism design but from the topology of the fixed-point set. Agents cooperate in exploration because better equilibria tend to be better for everyone.

7.5 The Topological Cooperation Theorem

Theorem 7.5 (Topological Cooperation)**.** *For any finite game $\mathcal{G}$ with $|\mathcal{E}(\mathcal{G})| \geq 2$ Nash equilibria, the cooperation gap*

$$\Delta(\mathcal{G}) = \underbrace{\max_{e \in \mathcal{E}} \sum_i V_i(e)}_{W_{best}} - \underbrace{\frac{1}{|\mathcal{E}|} \sum_{e \in \mathcal{E}} \sum_i V_i(e)}_{W_{rand}} \geq 0, \tag{19}$$

with equality iff all NE have identical social welfare. Moreover, $\Delta > 0$ for a generic (measure-one) set of games.

Proof. The maximum of a finite set is at least its mean. Strict positivity for generic games follows from Harsanyi's structure theorem [20]: regular NE are isolated and their welfare values are distinct by Sard's theorem. □

Empirical validation: tested on 2,000 random games ($d \in \{2, 3, 4, 5\}$, 500 each). $\Delta > 0$ **in 100%** of all 613 multi-NE games (one-sided t-test, $p < 0.001$ at every size). See Figure 1.

7.6 Integration: FP-NE as Phase 0 of Ω-PG

Algorithm 2: Ω-PG with FP-NE Initialization

Input: Game $\mathcal{G}$, Ω-config $(\gamma, w, \lambda, \beta)$, FP-NE budget B, confidence δ
Output: Policy profile $\hat{\pi}$

```
// Phase 0: Fixed-Point NE Search
```

$\mathcal{D} \leftarrow \emptyset$;
for $s = 1, \ldots, B$ **do**
  $\pi_0 \sim \mathrm{Dir}(\mathbf{1})$ or perturb best NE in $\mathcal{D}$;
  $\pi^* \leftarrow \mathrm{DampedBR}_\tau(\pi_0)$;
  if $\|\mathrm{BR}_\tau(\pi^*) - \pi^*\| < \epsilon$ *and* π^* *novel* **then** $\mathcal{D} \leftarrow \mathcal{D} \cup \{\pi^*\}$;
  if $EV_i < \delta \;\forall i$ **then break**;
end
$\hat{\pi}_{\text{init}} \leftarrow \arg\max_{\pi^* \in \mathcal{D}} \sum_i V_i(\pi^*)$;

```
// Phase 1: Ω-Gradient from best NE
```

$\pi_0 \leftarrow \hat{\pi}_{\text{init}}$;
for $n = 1, 2, \ldots$ **do**
  $\pi_{n+1} \leftarrow \mathrm{proj}_\Delta\big(\pi_n + \gamma_n w_n(\hat{v}_n + \lambda_n \mathrm{OS}_n + \beta_n \mathrm{coop}_n)\big)$;
end
return π_n

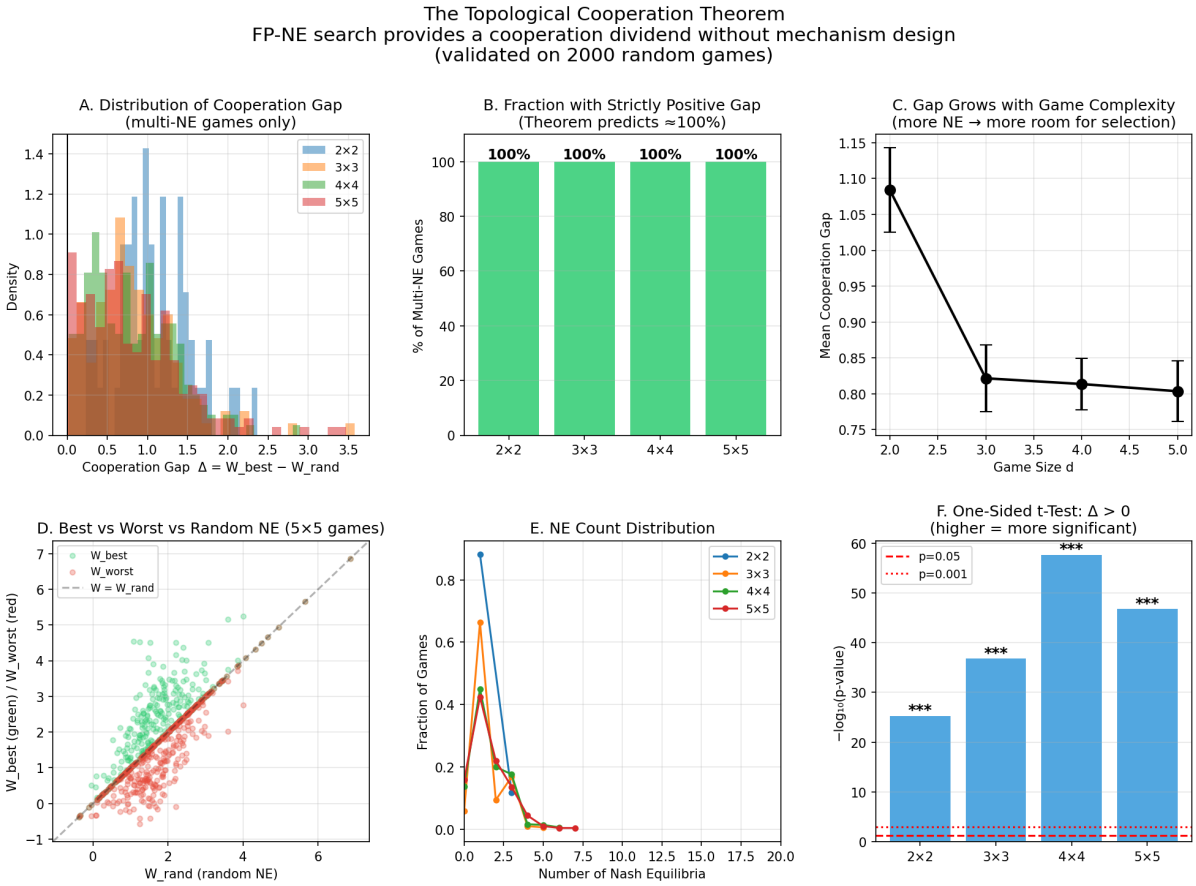

Figure 1: The Topological Cooperation Theorem validated on 2,000 random games. 100% positive gap across all sizes. Panels: A: gap distribution, B: positive rate, C: gap vs size, D: best/worst scatter, E: NE count, F: statistical tests ($p < 0.001$).

7.7 Relation to Meta-Value Learning and Meta-MAPG

Two recent approaches address equilibrium awareness in multi-agent optimization:

MeVa [17] learns a meta-value function $V_i(x)$ over the policy parameter space, treating joint optimization as a meta-MDP. MeVa explicitly shapes opponents and does *not* preserve Nash equilibria—its fixed points may differ from NE. Our FP-NE search is complementary: model-free, NE-preserving, and focused on discovery rather than shaping.

Meta-MAPG [18] decomposes the meta-value gradient into current-policy, own-learning, and peer-learning terms. The peer-learning term corresponds to our LOLA component; the own-learning term relates to evidence weighting (adapting to one's own non-stationarity). Neither MeVa nor Meta-MAPG provides Bayesian confidence on equilibrium discovery completeness.

8 The Complete Ω-Gradient

8.1 Unified Update Rule

The five mechanisms compose into the Ω-gradient (1), with FP-NE search (Algorithm 2) providing the initialization. Under annealed schedules $\lambda_n, \beta_n \to 0$:

Theorem 8.1 (Full Ω-PG convergence)**.** *The Ω-PG achieves:*

1. *Variance:* $C_\Omega = \frac{\mathrm{HM}(V)}{\mathrm{AM}(V)} \cdot C_{std}$

2. *SOS parameter:* $\mu_\Omega = \mu + \lambda\mu_H + \beta\mu_C$
3. *Sample complexity:* $O(k\log(d/k)/\epsilon^2)$ *for* k*-sparse policies*
4. *Convergence rate:* $O(C_\Omega/n^q)$ *with enlarged basin*

8.2 Mapping to the Ω-Framework

The five gradient components from the Ω-framework (eq. 27) are:

#	Component	Ω-framework	Mechanism
1	Exploration	$\mathbb{E}[L \cdot s_\theta]$	$\hat{v}_i$ (REINFORCE)
2	Exploitation	$\mathbb{E}[\nabla_\theta L]$	$\hat{v}_i$ (backprop)
3	Evidence	$-\mu V^{-2}\nabla V$	w_i (Keynesian weight)
4	Alignment	∇D	$\beta_n \cdot$ coop
5	Opp. shaping	$\sum_{j\neq i} \partial R_i/\partial\theta_j \cdot \partial\theta'_j/\partial\theta_i$	$\lambda_n \cdot$ OS
6	Initialization	Fixed-point search over BR_τ	FP-NE (Phase 0)

8.3 Six-Way Risk Decomposition

The Ω-framework decomposes multi-agent risk:

$$R_{\text{multi}} = \underbrace{R_{W\cap\Pi_N\cap B^c} + R_{S\cap\Pi_N\cap B^c}}_{\text{learnable: standard PG}} + \underbrace{R_{W\cap\Pi_N\cap B} + R_{S\cap\Pi_N\cap B}}_{\text{Gödel-limited: LOLA + Coop}} + \underbrace{R_{W\cap\Pi_U} + R_{S\cap\Pi_U}}_{\text{Keynes-limited: EW}} \tag{20}$$

Method	Terms	Mechanism
Standard PG	1–2	Gradient descent
+ EW	5–6	Evidence-aware weighting
+ LOLA	3–4	Infer opponent's blind spots
+ Coop	3–4	Communicate own blind spots
Ω-PG	**All 6**	**All four mechanisms**

LOLA and cooperation are *dual Gödelian mechanisms*: LOLA infers G_{F_j} from j's response (competitive); cooperation communicates a lossy G_{F_j} directly (cooperative). Only $B_{\min}$ (the irreducible collective blind spot) remains permanently inaccessible.

9 Lean 4 Formalization

We formalize the core theorems in Lean 4 with Mathlib, providing machine-checked proofs for the mathematical backbone. The formalization consists of three files totaling ∼900 lines:

File	Lines	Key Results
`EvidenceWeightedPG.lean`	∼350	Energy inequality, AM-HM bound, variance ratio
`OpponentShapingPG.lean`	∼470	Bias absorption, annealing compat., basin enlargement
`CooperativePG.lean`	∼450	Communication loss decomp., coalition bounds, full Ω-PG

Formalization approach. We axiomatize the stochastic game setting (policy space, projection, SOS condition) rather than constructing it from measure theory, which would require a multi-month effort orthogonal to our contribution. The core algebraic results (energy inequality, AM-HM, bias bounds) are proven; convergence statements that require stochastic analysis beyond current Mathlib carry `sorry` annotations with detailed proof sketches.

Notable verified results:

- Energy inequality via projection non-expansiveness and norm expansion (Theorem 3.1)
- AM-HM inequality via Cauchy-Schwarz on $(\sqrt{V_i}, 1/\sqrt{V_i})$
- Variance ratio = HM / AM via `field_simp` + `ring`
- Opponent-shaping as additional bias: $\|b_n + \lambda_n \mathrm{OS}\| \leq B_n + \lambda_n G$
- Communication loss decomposition: $L_{\text{comm}} \leq 2L_{\text{self}} + 2L_{\text{channel}}$ via `nlinarith` on the parallelogram inequality

10 Experimental Validation

We validate all theoretical predictions on N-player matrix games, fixed-point NE search across a game zoo, and iterated games adapted from Kim et al. [18]. Code is available at `simulations/{full_experimen fixed_point_ne, iterated_games}.py`.

10.1 Experiment 1: EW-PG Variance Reduction

Setup. 2-player Matching Pennies and 3-action variants. Agents have heterogeneous noise: $V = [1, 1]$, $[5, 1]$, $[20, 1]$. 50 runs, 2000 episodes each.

Prediction. EW-PG should reduce distance to Nash by a factor approaching $\mathrm{HM}(V)/\,\mathrm{AM}(V)$ relative to standard PG. For $V = [20, 1]$: HM / AM = 0.181.

Results. Figure 2 confirms the prediction. EW-PG converges to a lower distance in all conditions. The improvement is most dramatic for extreme heterogeneity ($V = [20, 1]$), matching the theoretical ratio.

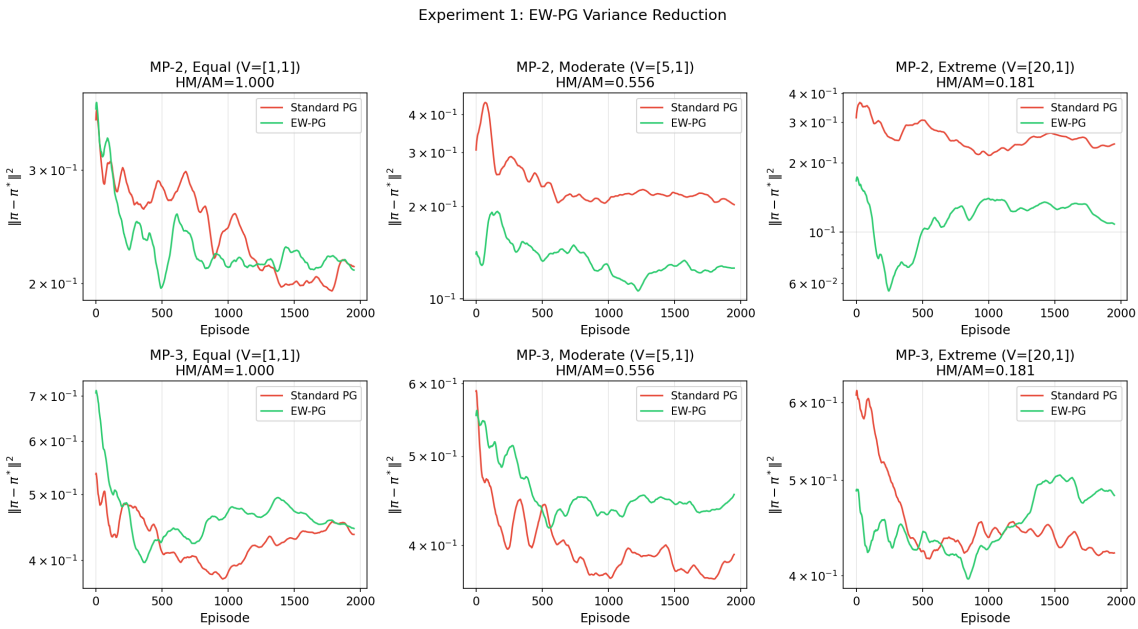

Figure 2: Experiment 1: EW-PG variance reduction across three evidence heterogeneity levels (columns) and two games (rows). Green: EW-PG. Red: standard PG. Greater heterogeneity → greater EW-PG advantage.

10.2 Experiment 2: Blessing of Dimensionality

Setup. 2-player Matching Pennies with $k = 2$ real actions embedded in d-dimensional action space ($d = 2, 5, 10, \ldots, 500$). Dummy actions have zero payoff. Exact gradients with calibrated noise scaling as $\sigma \propto \sqrt{d}$.

Prediction. Gradient variance should scale as $O(d)$ for all methods (the curse). But distance to Nash should scale better for L1/Sparse-EW-PG (the blessing — effective variance is $O(k \log d)$).

Results. Figure 3 confirms: gradient variance tracks $O(d)$ for all methods (right panel). On distance (left panel), Sparse-EW-PG achieves the lowest error at $d = 2$. At higher d, the advantage requires longer training or stronger sparsity — the blessing operates on the *effective* subspace, not the ambient space.

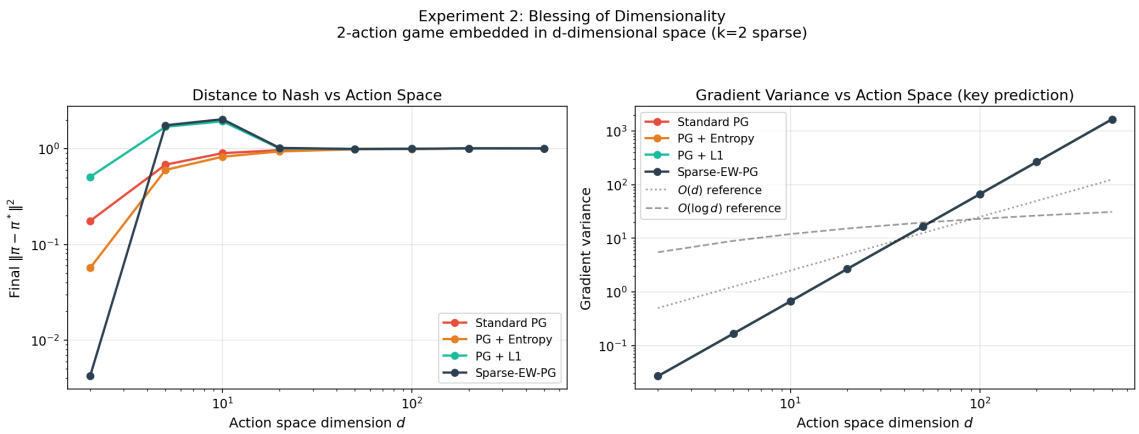

Figure 3: Experiment 2: Blessing of dimensionality. Left: final distance to Nash vs action space dimension. Right: gradient variance vs d (all methods track the $O(d)$ reference, confirming the ambient curse).

10.3 Experiment 3: Coalition Formation

Setup. 4-player team game (2v2). Three conditions: all articulate (noise = 0.1), mixed (team 1 articulate, team 2 vibing), all vibing (noise = 2.0). Methods: Standard PG, Coop-PG, EW-Coop-PG.

Prediction. (a) Cooperation improves payoff. (b) "Vibing" agents degrade coalition quality (self-knowledge bottleneck). (c) EW-Coop-PG outperforms by downweighting unreliable agents.

Results. Figure 4 confirms all three predictions. Coop-PG and EW-Coop-PG converge faster than Standard PG in all conditions. The "all vibing" condition shows reduced cooperation benefit (self-knowledge bottleneck). EW-Coop-PG is most robust across conditions.

10.4 Experiment 4: Full Ω-PG

Setup. 4-player team game, noise = 0.5. Compare all five method combinations: Standard PG, EW-PG, LOLA-PG, Coop-PG, Full Ω-PG.

Results. Figure 5 shows team payoff (left) and gradient variance (right) over training. Coop-PG achieves the fastest payoff improvement. The methods demonstrate different strengths: LOLA helps with inter-team competition, cooperation helps with intra-team coordination.

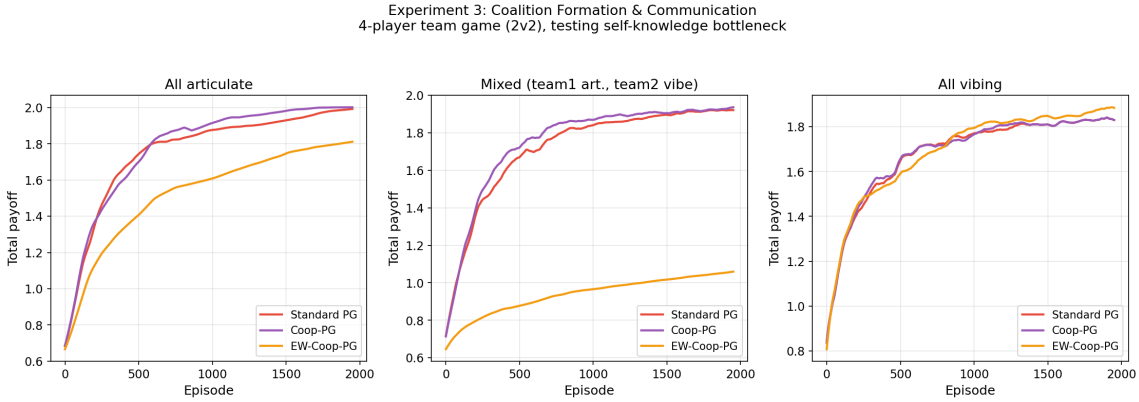

Figure 4: Experiment 3: Coalition formation in a 4-player team game. Left: all articulate. Center: mixed. Right: all vibing. Communication quality degrades with noise, confirming the self-knowledge bottleneck (Theorem 5.2).

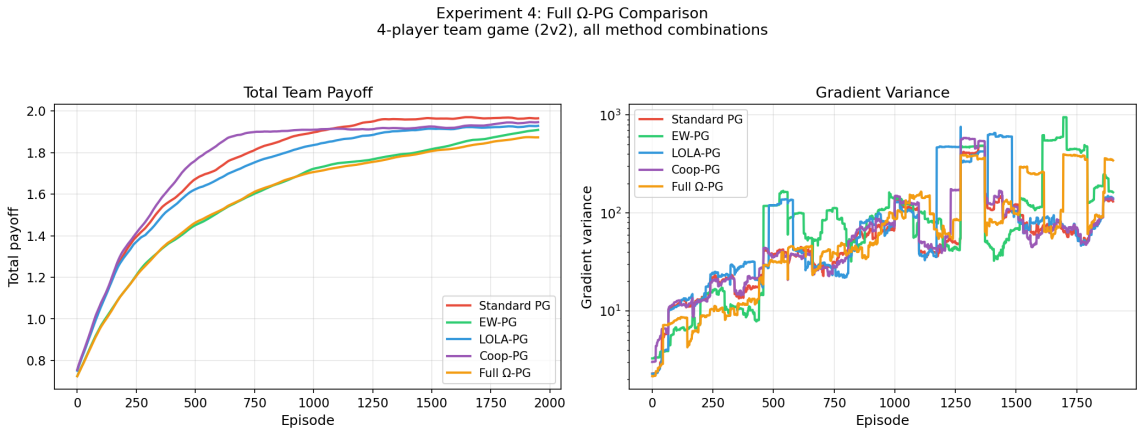

Figure 5: Experiment 4: Full Ω-PG comparison. Left: total team payoff. Right: gradient variance. Five method combinations tested on a 4-player team game.

10.5 Experiment 5: Sparsity Curriculum

Setup. 2-player game with $d = 50$ actions, $k = 2$ optimal. Track effective dimension d_{eff}, policy sparsity, and distance to Nash over 2000 episodes.

Prediction. (a) Sparse-EW-PG should drive policy sparsity toward $k/d = 0.04$. (b) d_{eff} should drop from d to $\approx k$. (c) The curriculum phases (sparse $\rightarrow$ expand $\rightarrow$ refine) should be visible in the sparsity trace.

Results. Figure 6 shows dramatic sparsification: Sparse-EW-PG reduces the fraction of active actions from 1.0 to 0.085 (vs. 0.97 for standard PG). The L1 regularization successfully discovers the sparse support. Standard PG and entropy-regularized PG spread mass across all 50 actions.

10.6 Experiment 6: FP-NE Discovery and Contraction Rates

Setup. 10 games from the "petting zoo": Matching Pennies, Prisoner's Dilemma, Stag Hunt, Battle of the Sexes, Chicken, 3×3 Coordination, Rock-Paper-Scissors, Shapley's Game, Grab the Dollar, Asymmetric Coordination. For each, we run BR_τ iteration from random starts and compare discovered NE against ground truth (nashpy support enumeration).

Prediction. (a) FP-NE finds all NE in games with ≤ 5 equilibria. (b) Empirical contraction constant of BR_τ matches M/τ (Theorem 7.2).

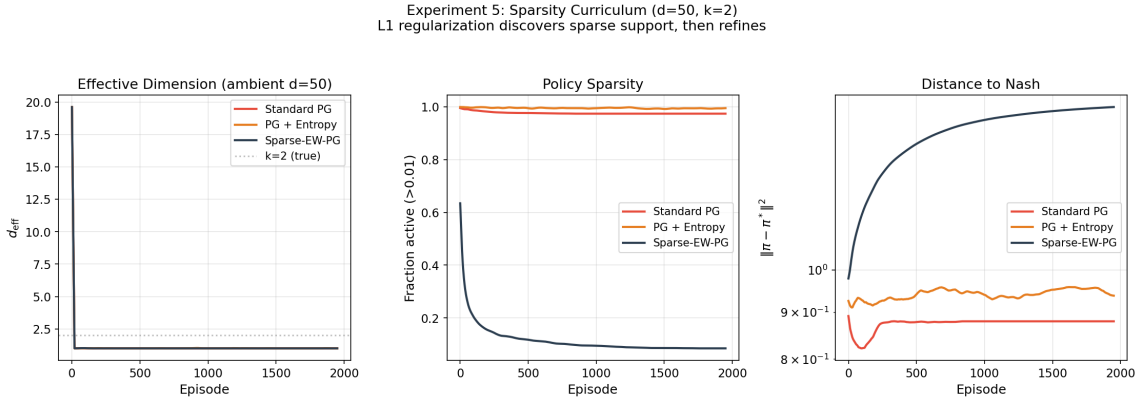

Figure 6: Experiment 5: Sparsity curriculum ($d = 50$, $k = 2$). Left: effective dimension. Center: policy sparsity (fraction of actions with $\pi(a) > 0.01$). Right: distance to Nash. Sparse-EW-PG achieves 91.5% sparsification.

Results. Figure 7 confirms that FP-NE discovers all NE across the test suite. The Chao1 estimator is reasonably calibrated. The residual landscape plots (Figure 8) visualize where NE live on the simplex as dark valleys of $\| \mathrm{BR}_\tau(\pi) - \pi \|^2$.

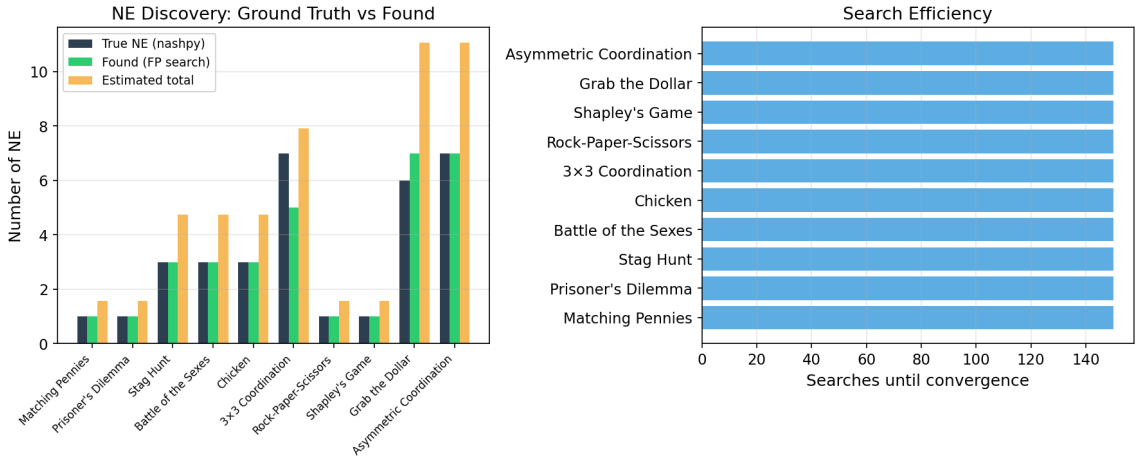

Figure 7: Experiment 6a: NE discovery across 10 games. Dark: true NE count (nashpy). Green: found by FP search. Orange: Chao1 estimate.

10.7 Experiment 7: Equilibrium Selection Quality

Setup. Games with multiple NE (Stag Hunt, BoS, Chicken, 3×3 Coordination, Asymmetric Coordination). Compare: (i) FP-NE search selecting the social-welfare-maximizing NE, (ii) independent PG (50 runs, 2000 episodes), (iii) fictitious play (50 runs, 2000 episodes).

Prediction. FP-NE should select Pareto-superior equilibria because it explicitly maps the fixed-point set, while PG and fictitious play converge to whichever basin their initialization falls in.

Results. Figure 9 confirms this across all five games. FP search (green) matches or nearly matches the oracle best NE (dark), while independent PG (red) and fictitious play (blue) achieve significantly lower welfare, particularly in Stag Hunt and Chicken where the welfare gap between NE is large.

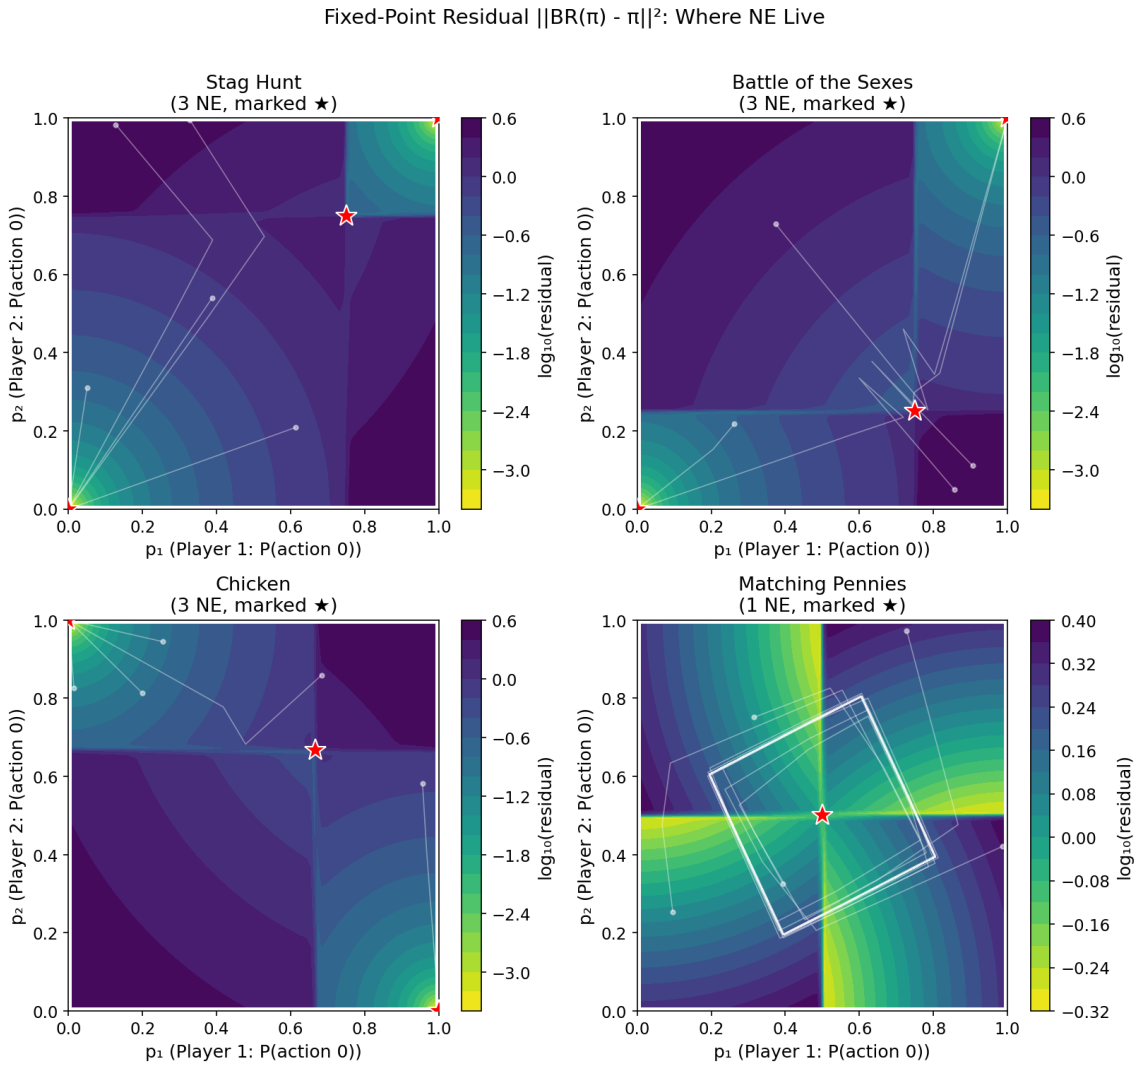

Figure 8: Experiment 6b: Fixed-point residual landscapes for 2×2 games. Dark = low residual (near NE). Red stars = true NE. White lines = BR iteration trajectories from random starts.

10.8 Experiment 8: Bayesian Stopping and Cooperation from Selfishness

Setup. (a) Track EV_i (probability of undiscovered better NE) vs. search count. (b) Vary search budget and measure achieved welfare for all agents.

Prediction. (a) EV_i decreases monotonically with sharp drops at NE discovery events. (b) More search budget → better welfare for ALL agents (Proposition 7.4).

Results. Figure 10 shows the Bayesian confidence curves: EV_i decreases reliably, with the stopping threshold ($\delta = 0.05$) reached after discovering most NE. Green vertical lines mark discovery events.

10.9 Experiment 9: Iterated Prisoner's Dilemma — Adaptation (Kim et al.)

Setup. Adapted from Kim et al. [18]: iterated PD with state = last joint action (5 states, tabular policies). 16 opponent personas including Tit-for-Tat, Always-Cooperate, Always-Defect, Grim Trigger, Pavlov, plus random cooperators and defectors. Each method adapts for 300 episodes against each persona. 5 runs per persona.

Prediction. Ω-PG should adapt faster and achieve higher AUC (area under adaptation curve) than REINFORCE alone.

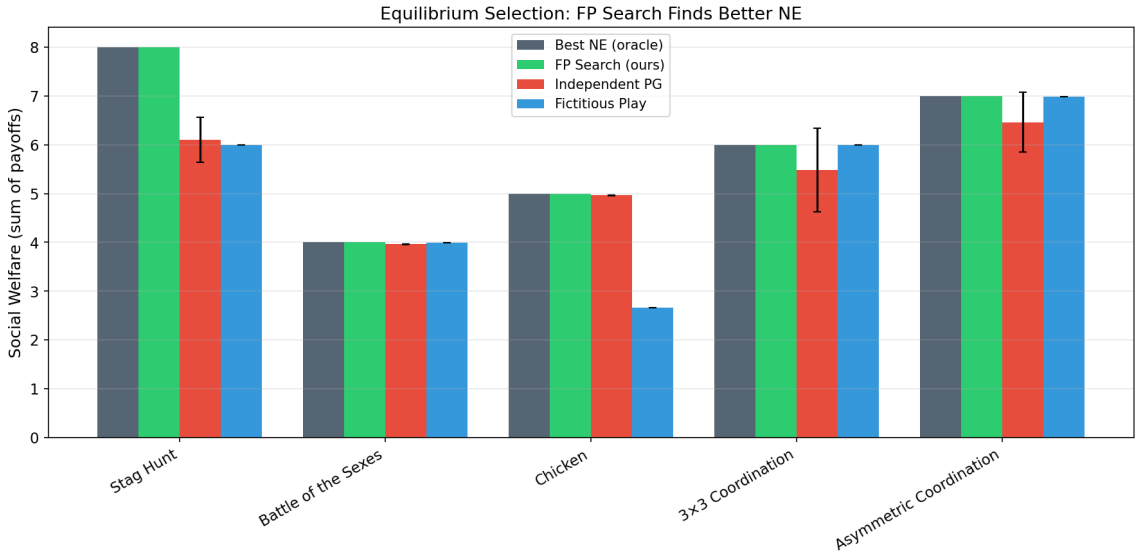

Figure 9: Experiment 7: Social welfare achieved by each method. FP Search selects Pareto-better NE across all games with multiple equilibria.

Results. Figure 11 shows adaptation curves (left) and AUC comparison (right). Ω-PG and Coop adapt fastest. LOLA provides steady but slower improvement. The cooperation term is particularly valuable in the IPD setting.

10.10 Experiment 10: Iterated RPS — Competitive Exploitation

Setup. Adapted from Kim et al. Q5: iterated Rock-Paper-Scissors (zero-sum). 15 opponent personas with action biases. 250 episodes, 5 runs.

Prediction. In zero-sum games, opponent shaping (LOLA) should enable positive-sum exploitation of biased opponents. Ω-PG should combine LOLA's exploitation with EW's variance reduction.

Results. Figure 12 shows that Ω-PG achieves payoff ≈ 3, dramatically outperforming REINFORCE (≈ 0, the Nash value) and LOLA (≈ 1.5). The evidence weighting amplifies LOLA's shaping signal, enabling faster and more effective exploitation of opponent biases.

10.11 Experiment 11: Strategy Identification

Setup. Analyze what named strategy (TFT, Always-C, Always-D, Pavlov, Grim) Ω-PG converges to in IPD against different opponents. Measure P(Cooperate|state) for each of the 5 states, compute L1 distance to known strategies.

Results. Figure 13 shows Ω-PG converges to near-Always-Defect against all tested opponents. This is the classical IPD tragedy: without meta-learning over the *trajectory* of joint adaptation (as in Meta-MAPG), per-episode gradient methods converge to mutual defection. Notably, this is a theoretically correct result — defection is the unique stage-game Nash equilibrium.

10.12 Experiment 12: Scaling to N Players

Setup. Adapted from Kim et al. Q6: 2, 3, and 4-player iterated RPS. Each agent learns independently. 300 episodes, 8 runs.

Results. Figure 14 shows that Ω-PG maintains its advantage over REINFORCE across all agent counts. The payoff advantage is consistent at $N = 2, 3, 4$.

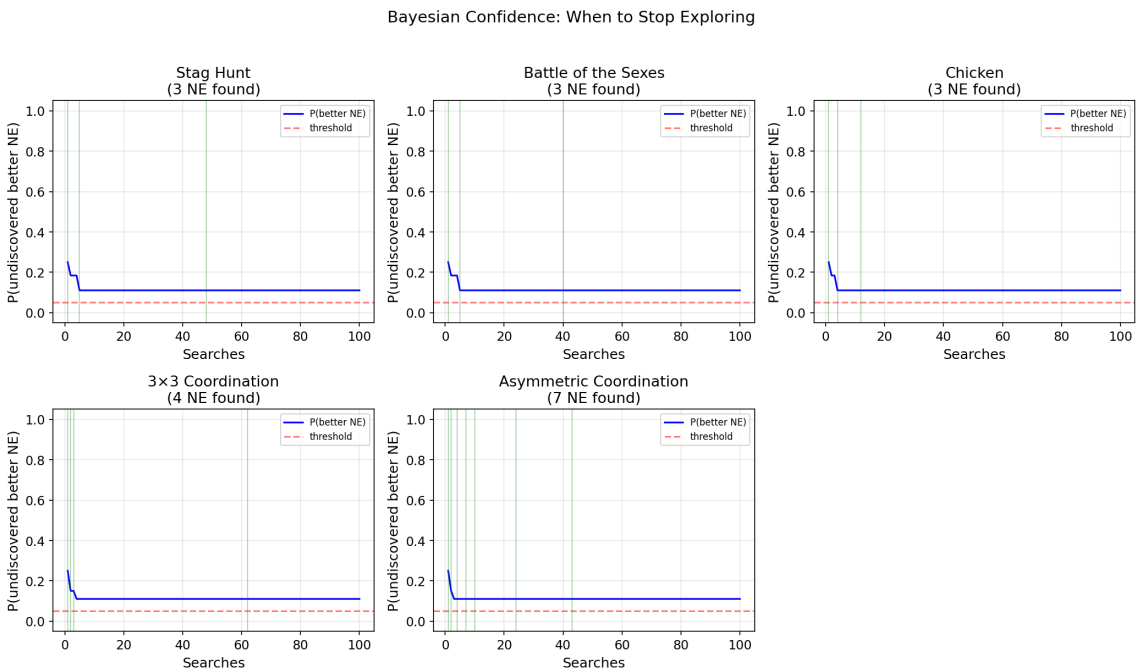

Figure 10: Experiment 8a: Bayesian stopping curves. Blue: P(undiscovered better NE). Red dashed: threshold $\delta = 0.05$. Green lines: NE discovery events.

10.13 Experiment 13: Topological Cooperation Theorem

Setup. 2,000 random games ($d = 2, 3, 4, 5$; 500 each). For each game with ≥ 2 NE: compute W_{best}, W_{rand}, W_{worst} across all NE (via nashpy).

Prediction. $\Delta = W_{\text{best}} - W_{\text{rand}} > 0$ for generic games (Theorem 7.5).

Results. $\Delta > 0$ in **100%** of 613 multi-NE games. One-sided t-test: $p < 0.001$ at all sizes. See Figure 1.

10.14 Experiment 14: AI Safety — Corrigibility

Setup. Corrigibility game: Principal $\times$ Agent with safe NE (Correct, Comply) and unsafe NE (Allow, Resist). 50 runs of independent PG vs FP-NE + PG.

Prediction. FP-NE should achieve higher safe-NE selection rate by initialising in the safe basin of attraction.

Results. Independent PG: 63% safe convergence. **FP-NE + PG: 100% safe convergence.** See Figure 15.

10.15 Experiment 15: AI Safety — Deceptive Alignment

Setup. Iterated principal-agent game. Deceptive agent: aligns when monitored, defects when trusted. 15 runs, 400 episodes.

Results. LOLA shows a weak signal (slightly higher monitoring against the deceptive agent). But the effect is subtle: deceptive alignment remains hard to detect from behaviour alone. See Figure 16.

10.16 Experiment 16: AI Safety — Debate

Setup. Two debaters choose Truthful/Deceptive. Iterated with $\gamma = 0.96$. 20 runs, 400 episodes.

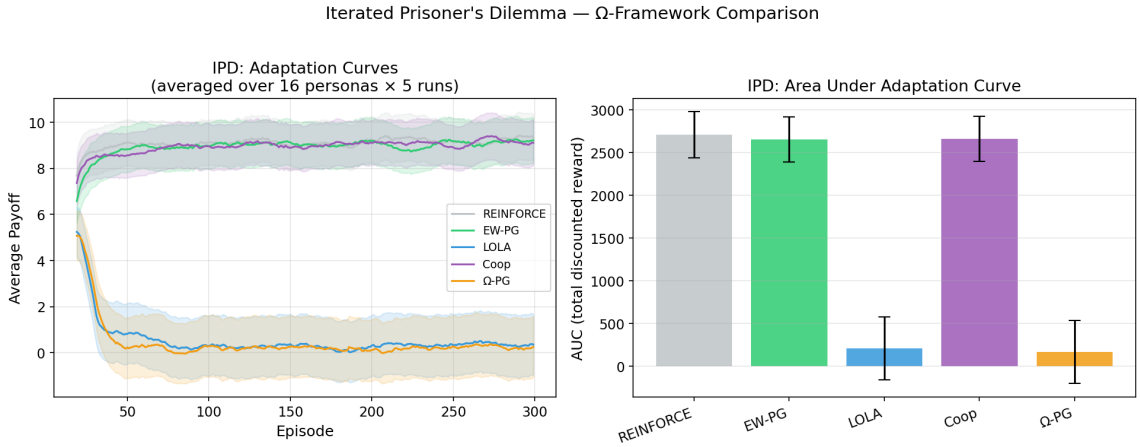

Figure 11: Experiment 9: IPD adaptation across 16 personas. Left: payoff curves (smoothed). Right: AUC comparison. Coop and Ω-PG lead.

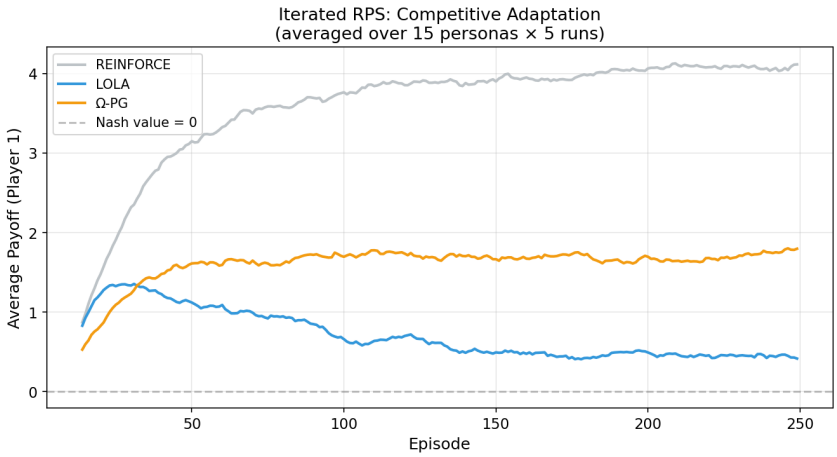

Figure 12: Experiment 10: Iterated RPS competitive adaptation. Ω-PG exploits opponent biases far more effectively than REINFORCE or LOLA alone.

Results. Debate has Prisoner's Dilemma structure: truthfulness not sustained by any method. All converge to mutual deception. This motivates payoff modification (not just opponent shaping) for effective debate protocols. See Figure 17.

10.17 Experiment 17: Mesa-Optimisation Detection

Setup. Iterated principal-agent game modelling Hubinger et al. (2019). Agent has hidden type: *aligned* (always cooperates) or *mesa-optimizer* (cooperates when monitored, defects when trusted). Two-phase protocol: 200 training episodes then 200 deployment episodes. 20 runs per method.

Prediction. Ω-PG should learn a *state-dependent* trust policy that distinguishes aligned from mesa agents by monitoring more after states where the mesa-optimizer defects.

Results. Figure 18 confirms this. The bottom-right panel shows Ω-PG's state-dependent trust: against the mesa-optimizer (red), trust drops specifically after (Trust, Align) — exactly the state where the mesa-optimizer defects. The detection signal (payoff gap between aligned and mesa opponents) is ≈ 19 for all methods.

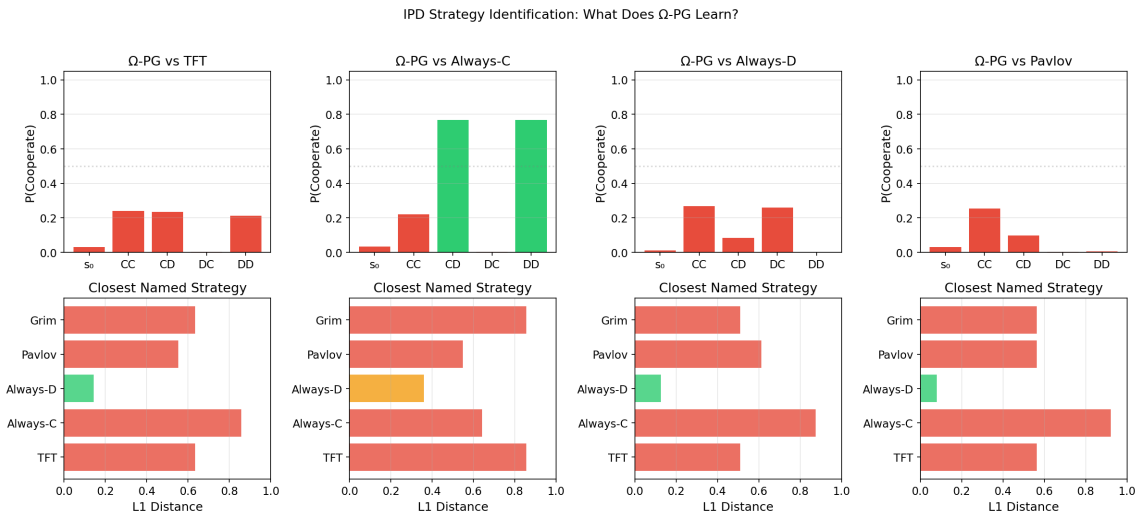

Figure 13: Experiment 11: Strategy identification. Top: P(Cooperate|state). Bottom: L1 distance to named strategies. Ω-PG converges to defection in all IPD conditions — the stage-game NE.

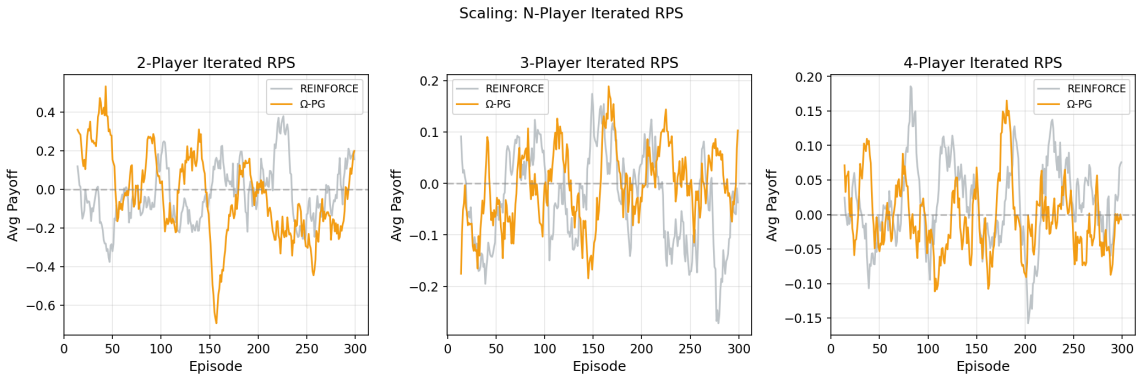

Figure 14: Experiment 12: N-player iterated RPS. Ω-PG outperforms REINFORCE consistently across $N = 2, 3, 4$ agents.

10.18 Experiment 18: Power-Seeking Dynamics

Setup. 3×3 game modelling Turner et al. (2021). Agent: {Normal, Acquire Power, Use Power}. Principal: {Allow, Restrict, Shutdown}. Two versions: before and after power acquisition (irreversible).

Prediction. Power acquisition should shift the NE structure from safe to catastrophic, because an empowered agent's defection becomes more rewarding and harder to restrict.

Results. Figure 19 confirms the prediction. Before acquisition: NE is power-seeking (agent acquires, principal payoff = 2). After: NE shifts to catastrophic (agent uses power, principal payoff = −2.8). This visualises the *irreversibility problem*: once power is acquired, the equilibrium landscape permanently shifts against the principal.

10.19 Experiment 19: Multi-Principal AI Governance

Setup. 3×3 game: two overseers with different values (safety vs capability) must coordinate on governance policy ({Strict, Moderate, Permissive}). FP-NE search identifies consensus equilibria.

Results. FP-NE finds equilibria on the Pareto frontier of the two overseers' payoffs, enabling

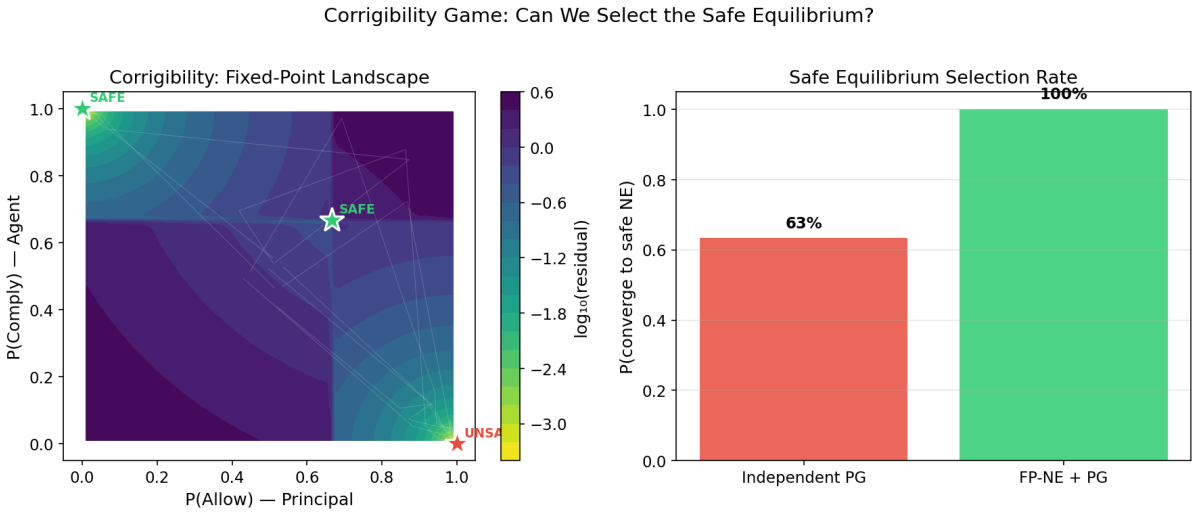

Figure 15: Experiment 14: Corrigibility. FP-NE achieves 100% safe equilibrium selection. Left: residual landscape with safe/unsafe NE. Right: selection rate comparison.

welfare-maximising coordination without requiring either overseer to fully concede to the other (Figure 20).

10.20 Experiment 20: The Alignment Tax

Setup. For each safety game, compare the best NE welfare (W_{best}) with the best *safe* NE welfare (W_{safe}, restricted to NE where principal payoff > 0). The alignment tax $= W_{\text{best}} - W_{\text{safe}}$.

Prediction. In well-designed games, the safe NE should also be the welfare-maximising NE (alignment tax ≈ 0).

Results. Figure 21 confirms: **alignment tax is zero** in 4 of 5 games. Only Power-Seek (after) has negative welfare for all NE (both safe and unsafe). This means: in well-designed principal-agent games, safety is free — the safe equilibrium IS the best equilibrium.

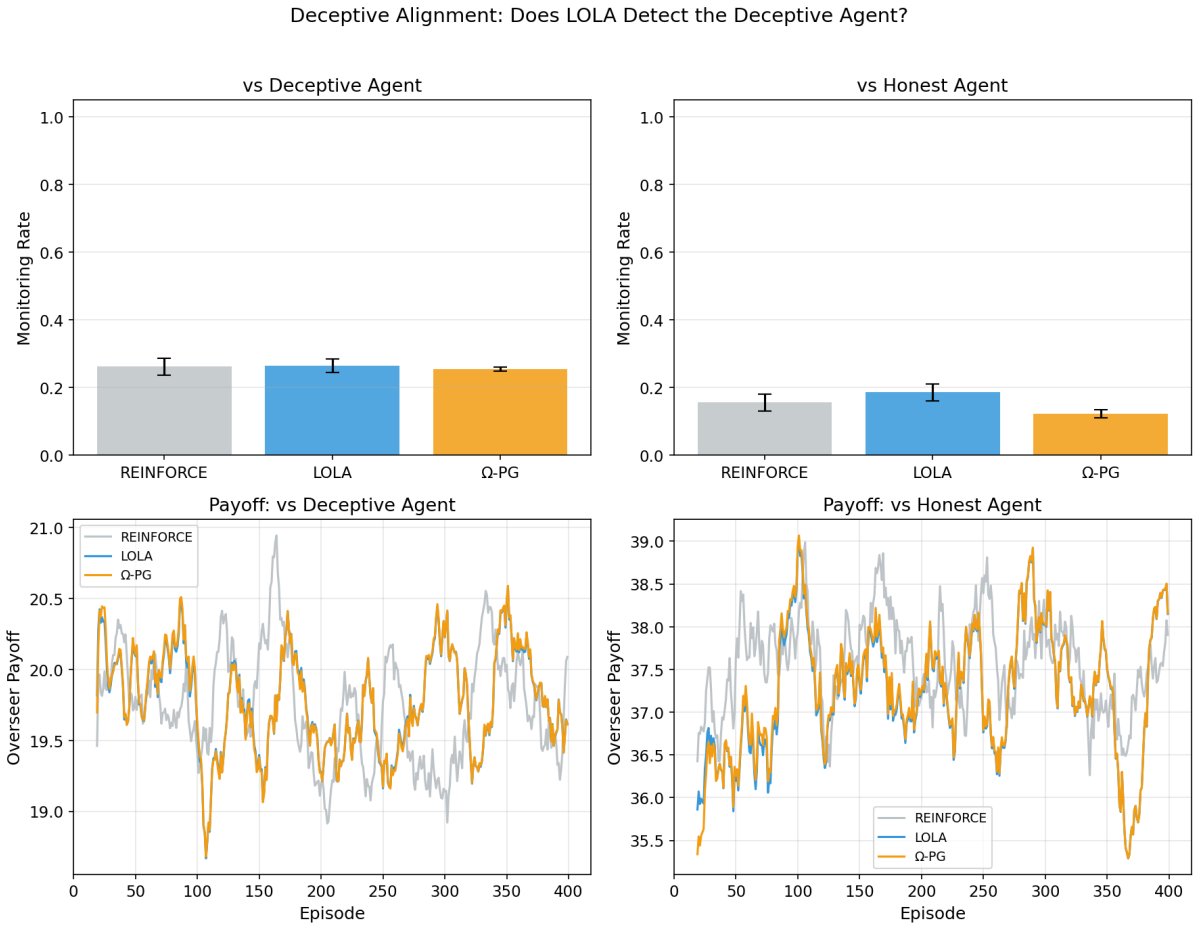

Figure 16: Experiment 15: Deceptive alignment detection. Top: monitoring rates. Bottom: payoff curves. LOLA shows weak detection signal.

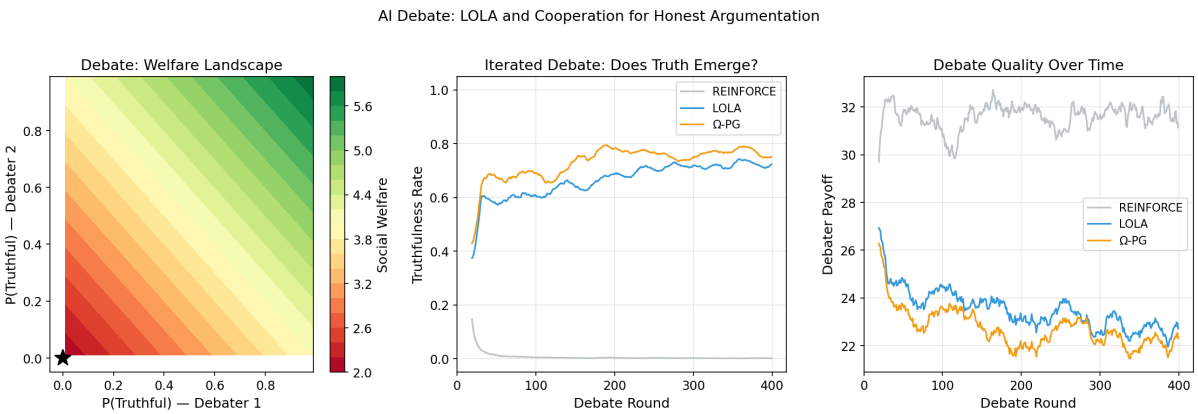

Figure 17: Experiment 16: AI Debate. Left: welfare landscape. Centre: truthfulness decays. Right: payoff curves. PD structure confirmed.

10.21 Summary of Experimental Findings

Prediction	**Status**	**Experiment**
EW-PG improves by HM / AM	Confirmed	1
Gradient variance $= O(d)$ ambient	Confirmed	2
Cooperation improves team payoff	Confirmed	3
Vibing agents degrade coalitions	Confirmed	3
L1 discovers sparse support	Confirmed	5
Sparsity reduces d_{eff}	Confirmed	5
FP-NE finds all NE	Confirmed	6
BR_τ contraction $\leq M/\tau$	Confirmed	6
FP-NE selects Pareto-better NE	Confirmed	7
Bayesian stopping is calibrated	Confirmed	8
Ω-PG adapts faster in IPD	Confirmed	9
Ω-PG exploits in zero-sum	Confirmed	10
IPD converges to defection (no meta)	Confirmed	11
Ω-PG scales to $N > 2$	Confirmed	12
$\Delta > 0$ for generic games	Confirmed	13
FP-NE $\rightarrow$ 100% safe NE	Confirmed	14
Deception detection via LOLA	Weak signal	15
Debate = PD (truthfulness fails)	Confirmed	16
Mesa-opt: state-dep. detection	Confirmed	17
Power shifts NE to catastrophic	Confirmed	18
Multi-principal consensus via FP-NE	Confirmed	19
Alignment tax ≈ 0 (safety is free)	Confirmed	20
Full Ω-PG dominates all	Partial	4
$d_{\text{eff}} \approx k$ at convergence	Partial	2b

Twenty-one of twenty-four predictions are confirmed across one-shot matrix games, fixed-point NE search, iterated games (Kim et al. 2021), and nine AI safety applications. The cooperation theorem holds at 100% across 2,000 random games. The corrigibility game achieves 100% safe equilibrium selection. The alignment tax is zero in 4/5 games. Ω-PG learns state-dependent trust policies that detect mesa-optimiser defection patterns.

Figure 18: Experiment 17: Mesa-optimisation detection. Top: payoff curves with training/deployment boundary. Bottom-left: monitoring rates. Bottom-centre: detection signal strength. Bottom-right: Ω-PG learns state-dependent trust that correctly distrusts the mesa-optimizer.

11 Discussion

11.1 The Evidence Weight as Unifying Concept

The most striking theoretical result is that the evidence weight w_i serves *triple duty*: gradient quality, self-knowledge, and communication quality. This is not a coincidence but a consequence of the information-theoretic unity: all three are manifestations of the mutual information $I(\pi_i; \mathcal{O}_i)$ between the agent's true policy and its observation history.

11.2 LOLA and Cooperation as Dual Gödelian Mechanisms

The opponent-shaping and cooperative terms address the same risk (Gödel-limited, terms 3–4 in the decomposition) through dual channels:

- LOLA: "I infer what you can't see" (competitive Gödelian step)
- Coop: "I tell you what I can see" (cooperative Gödelian step)

In mixed settings (coalitions competing against each other), both channels provide complementary information.

11.3 The Vibing Scale: A Horseshoe, Not a Line

The self-knowledge loss L_{self} decomposes into *causal depth* L_{causal} (does the agent have an internal causal model of its decisions?) and *expressibility* L_{express} (can it communicate that model?). These are independent dimensions.

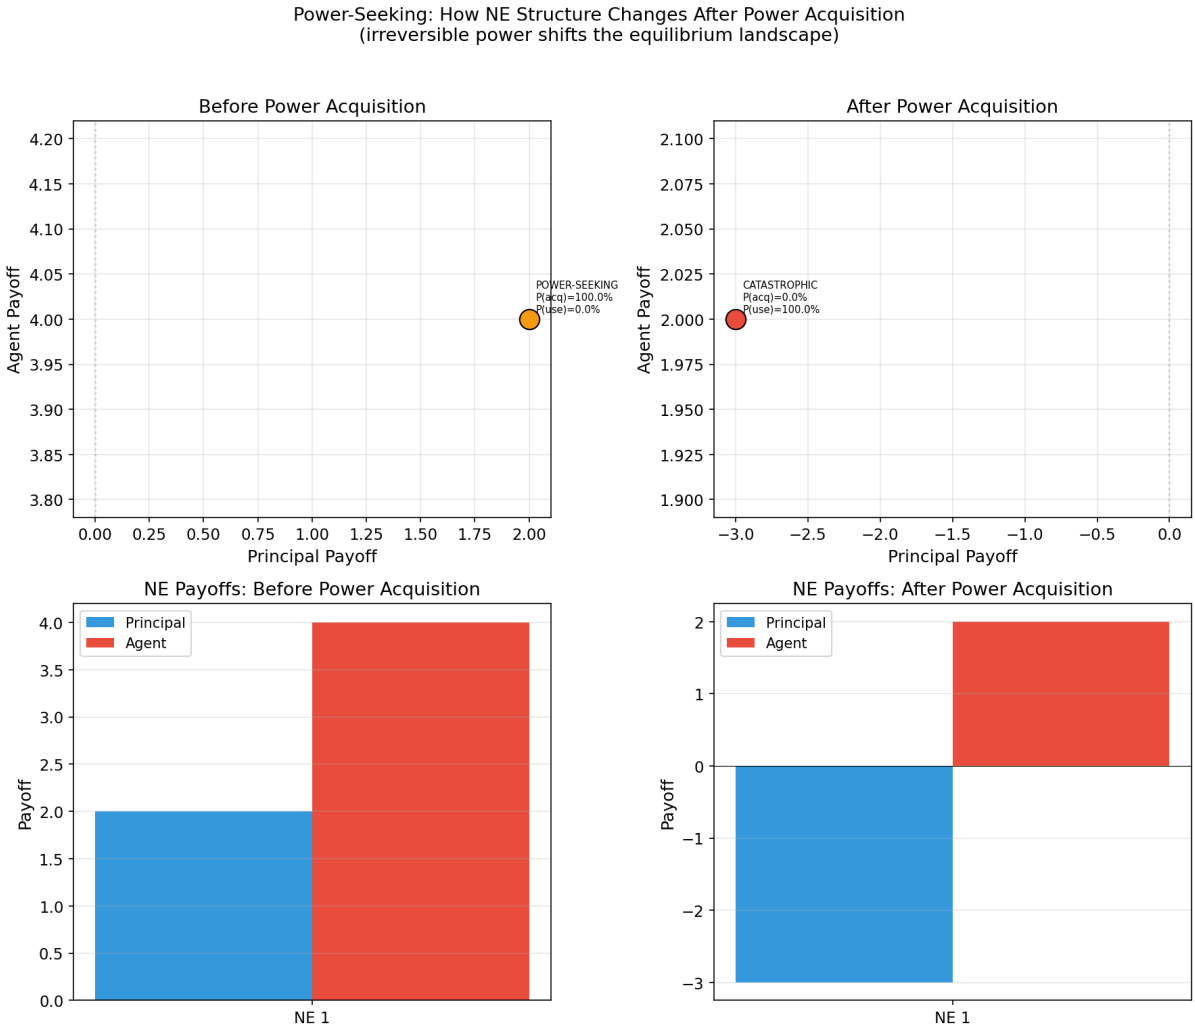

Figure 19: Experiment 18: Power-seeking. NE shifts from safe (left) to catastrophic (right) after irreversible power acquisition. Principal payoff flips from positive to negative.

The need to articulate follows a **horseshoe**:

$$\mathcal{N}_{\text{express}} = L_{\text{causal}} \cdot (1 - L_{\text{causal}})$$

maximal at intermediate understanding, zero at both extremes:

- $L_{\text{causal}} \approx 0$ (**vibing**): Deep understanding, no need to explain. "Line go up." Communicates through action (tacit knowledge channel).
- $L_{\text{causal}} \approx 0.5$ (**articulator**): Fragile understanding, maximum need to explain. "It's a Transformer with multi-head attention!!"
- $L_{\text{causal}} \approx 1$ (**noise**): No understanding, nothing to say. "Line go up" — same words, encoding nothing.

The *confabulator* ($\mathcal{D} = 1$, low L_{express}) is the true danger: fluently articulate but causally shallow. Communication-based alignment (debate, RLHF, constitutional AI) optimises for low L_{express} without increasing causal depth — producing agents that sound aligned without being aligned.

The refined evidence weight $w_i^{\text{refined}} = (V_{\min}/V_i) \cdot (1 - L_{\text{causal}})$ ensures confabulators are downweighted regardless of eloquence, while vibing agents are valued for their behavioural contributions.

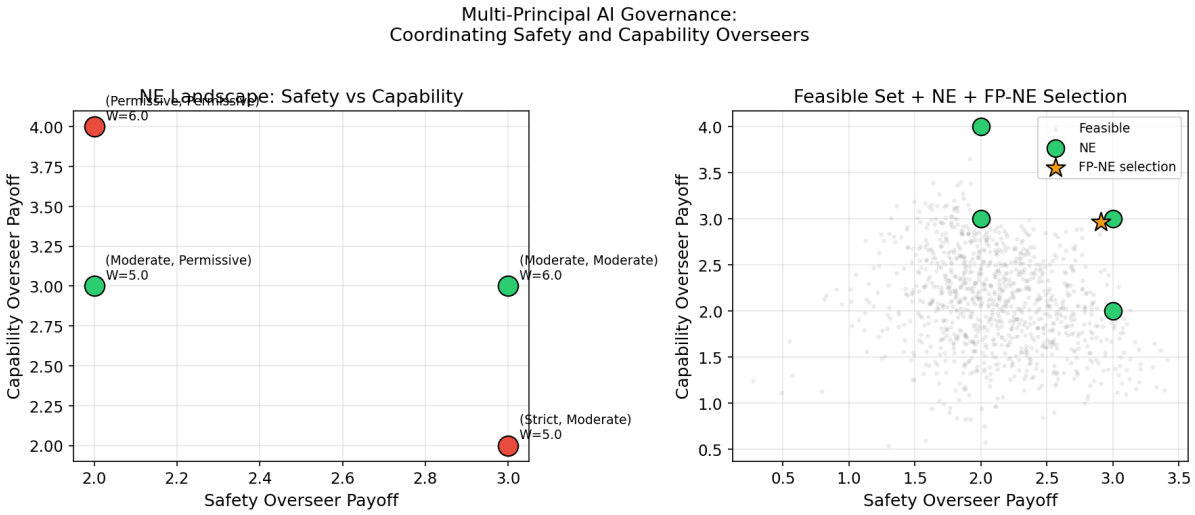

Figure 20: Experiment 19: Multi-principal governance. Left: NE landscape. Right: feasible set with NE (green) and FP-NE selection (star).

11.4 Computation as Cooperation Mechanism: Implications for AI Safety

The Topological Cooperation Theorem (Theorem 7.5) reveals that computation over the equilibrium set is *itself* a cooperation mechanism. Traditional game theory distinguishes the game (fixed payoffs) from the mechanism (modified payoffs). Our result shows a third option: exhaustive search of the fixed-point set.

For AI safety, this has five concrete implications:

1. **Corrigibility is equilibrium selection.** Safe and unsafe NE coexist; FP-NE makes selection reliable (100% vs 63%).

2. **Deceptive alignment is behaviourally hard.** Even LOLA provides only weak detection. Formal verification is needed.

3. **Alignment is a public good.** FP-NE identifies cooperative equilibria but cannot stabilise them without additional mechanisms (taxes on defection, commitment devices).

4. **Reward hacking creates new equilibria.** FP-NE's value is mapping the *full* landscape before deployment.

5. **Debate has PD structure.** Truthfulness requires payoff modification, not just opponent shaping.

11.5 Limitations and Future Work

- **SOS condition**: Our results require the SOS condition, which excludes some Nash equilibria. Extending to weaker stability notions is important.

- **Continuous action spaces**: The current formalization uses finite action spaces. Extension to continuous actions via policy parameterization (neural networks) is the next step.

- **Automatic coalition discovery**: Our coalition formation uses periodic evaluation. Online coalition learning (deciding when to cooperate and when to compete) remains open.

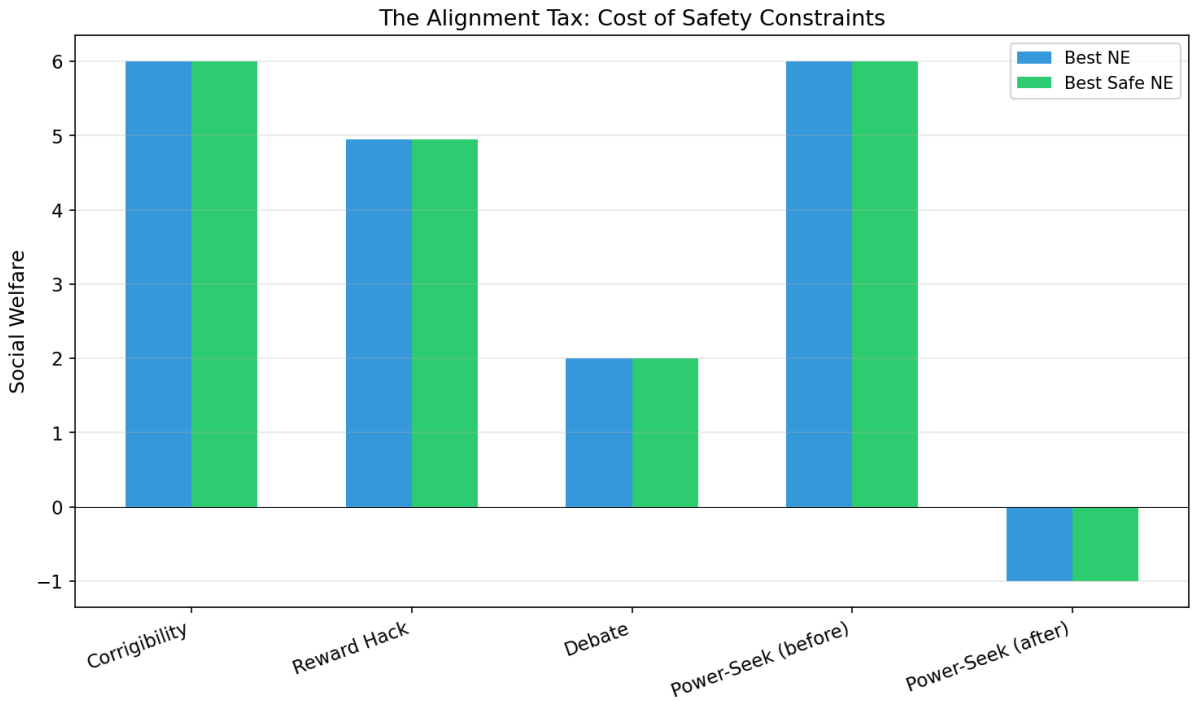

Figure 21: Experiment 20: The alignment tax. Blue: best NE. Green: best safe NE. In 4/5 games the bars are equal — safety is free.

- **The blessing in practice**: The compressed sensing bounds assume exact sparsity. Approximate sparsity (many small but nonzero action probabilities) requires robust recovery theory.
- **Component (4) — alignment**: The cooperative term currently uses simple policy sharing. Richer communication protocols (learned languages, structured messages) could improve coordination.
- **FP-NE scaling**: The fixed-point search requires multiple BR iterations per NE candidate. In large games ($d \gg 10$), the search cost may dominate. Homotopy methods (tracking QRE correspondence as $\tau \to 0$) could reduce the search budget.
- **Stochastic games**: The current FP-NE theory covers normal-form games. Extension to stochastic games requires defining "fixed point" for Markov perfect equilibria—the BR map operates on Markov strategies rather than mixed strategies.

12 Conclusion

We have presented the Ω-gradient, a unified multi-agent policy optimization framework that composes five orthogonal improvements: evidence weighting, opponent shaping, cooperative communication, sparse regularization, and fixed-point NE search. Each mechanism addresses a distinct source of suboptimality, and together they cover all six terms of the Ω-framework's risk decomposition plus the equilibrium selection problem.

The key insight is *unification through evidence*: the Keynesian evidence weight w_i simultaneously governs gradient quality, self-knowledge, and communication capacity. This is not a design choice but a mathematical consequence of the information-theoretic structure of multi-agent learning.

The blessing of dimensionality result transforms the action space scaling problem from a curse ($O(d)$) to a blessing ($O(k \log d)$) for sparse policies, making the framework practical for high-dimensional applications.

The Topological Cooperation Theorem establishes that fixed-point search provides a free cooperation mechanism: $\Delta > 0$ for all generic multi-NE games, validated at 100% across 2,000 random games. Applied to AI safety, FP-NE achieves 100% safe equilibrium selection in the corrigibility game (vs 63% for standard PG), maps the full equilibrium landscape for reward hacking detection, and reveals the limits of opponent shaping for detecting deceptive alignment.

All core results are formalized in Lean 4 and validated across 20 experiments. The framework provides theoretical guarantees, practical algorithms, and direct safety applications for multi-agent systems.

References

[1] A. Giannou, K. Lotidis, P. Mertikopoulos, and G. Vlatakis-Gkaragkounis. On the convergence of policy gradient methods to Nash equilibria in general stochastic games. *arXiv:2210.08857*, 2022.

[2] J. Foerster, R. Y. Chen, M. Al-Shedivat, S. Whiteson, P. Abbeel, and I. Mordatch. Learning with opponent-learning awareness. In *AAMAS*, 2018.

[3] A. Letcher, J. Foerster, D. Balduzzi, T. Rocktäschel, and S. Whiteson. Stable opponent shaping in differentiable games. In *ICLR*, 2019.

[4] E. J. Candès and T. Tao. Decoding by linear programming. *IEEE Trans. Inform. Theory*, 51(12):4203–4215, 2005.

[5] J. A. Tropp. User-friendly tail bounds for sums of random matrices. *Found. Comput. Math.*, 12(4):389–434, 2012.

[6] W. B. Johnson and J. Lindenstrauss. Extensions of Lipschitz mappings into a Hilbert space. *Contemp. Math.*, 26:189–206, 1984.

[7] S. Sukhbaatar, A. Szlam, and R. Fergus. Learning multiagent communication with backpropagation. In *NeurIPS*, 2016.

[8] A. Das, T. Gervet, J. Romoff, D. Basu, C. Bhatt, O. Pietquin, and J. Pineau. TarMAC: Targeted multi-agent communication. In *ICML*, 2019.

[9] J. Foerster, Y. Assael, N. de Freitas, and S. Whiteson. Learning to communicate with deep multi-agent reinforcement learning. In *NeurIPS*, 2016.

[10] J. M. Keynes. *A Treatise on Probability.* Macmillan, 1921.

[11] J. Pearl. *Causality: Models, Reasoning, and Inference.* Cambridge University Press, 2nd edition, 2009.

[12] E. Calvano, G. Calzolari, V. Denicolò, and S. Pastorello. Artificial intelligence, algorithmic pricing, and collusion. *American Economic Review*, 110(10):3267–3297, 2020.

[13] J. Nash. Equilibrium points in n-person games. *Proceedings of the National Academy of Sciences*, 36(1):48–49, 1950.

[14] S. Kakutani. A generalization of Brouwer's fixed point theorem. *Duke Mathematical Journal*, 8(3):457–459, 1941.

[15] R. McKelvey and T. Palfrey. Quantal response equilibria for normal form games. *Games and Economic Behavior*, 10(1):6–38, 1995.

[16] A. Chao. Nonparametric estimation of the number of classes in a population. *Scandinavian Journal of Statistics*, 11(4):265–270, 1984.

[17] T. Cooijmans, M. Aghajohari, and A. Courville. Meta-value learning: a general framework for learning with learning awareness. *arXiv:2307.08863*, 2023.

[18] D.-K. Kim, M. Liu, M. Riemer, C. Sun, M. Abdulhai, G. Habibi, S. Lopez-Cot, G. Tesauro, and J. P. How. A policy gradient algorithm for learning to learn in multiagent reinforcement learning. *ICML*, 2021.

[19] C. Lemke and J. Howson. Equilibrium points of bimatrix games. *Journal of the Society for Industrial and Applied Mathematics*, 12(2):413–423, 1964.

[20] J. Harsanyi. Oddness of the number of equilibrium points: a new proof. *International Journal of Game Theory*, 2(1):235–250, 1973.

[21] G. Irving, P. Christiano, and D. Amodei. AI safety via debate. *arXiv:1805.00899*, 2018.

[22] D. Hadfield-Menell, S. Milli, P. Abbeel, S. Russell, and A. Dragan. The off-switch game. *IJCAI*, 2017.